

Import Dependencies

```
In [1]: # Import standard packages for data manipulation and visualization
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Import tqdm to help visualize for loop iteration and time
from tqdm import tqdm

# Use missingno for easy missing data visualizations
import missingno as msno

# Import sklearn packages
from sklearn.ensemble import RandomForestClassifier
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.impute import SimpleImputer
from sklearn.compose import ColumnTransformer
from sklearn.model_selection import train_test_split, cross_val_score, GridSearchCV, Rep
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import OrdinalEncoder, OneHotEncoder, PowerTransformer, Func

from sklearn.metrics import classification_report, plot_confusion_matrix, accuracy_scor

import optuna
from optuna import visualization
import xgboost as xgb

import warnings
warnings.filterwarnings('ignore')
```

Read in data

```
In [2]: # Read in HR data
df_aug_train = pd.read_csv('aug_train.csv')
df_aug_test = pd.read_csv('aug_test.csv')

df = pd.concat( [df_aug_train, df_aug_test] )

# Observe details about dataset
df.info()

# Peek at data
df.head()
```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 21287 entries, 0 to 2128
Data columns (total 14 columns):

#	Column	Non-Null Count	Dtype
0	enrollee_id	21287 non-null	int64
1	city	21287 non-null	object
2	city_development_index	21287 non-null	float64
3	gender	16271 non-null	object

```

4  relevent_experience      21287 non-null object
5  enrolled_university      20870 non-null object
6  education_level          20775 non-null object
7  major_discipline         18162 non-null object
8  experience                21217 non-null object
9  company_size              14727 non-null object
10 company_type              14513 non-null object
11 last_new_job              20824 non-null object
12 training_hours           21287 non-null int64
13 target                   19158 non-null float64
dtypes: float64(2), int64(2), object(10)
memory usage: 2.4+ MB

```

Out[2]:

	enrollee_id	city	city_development_index	gender	relevent_experience	enrolled_university	education_level
0	8949	city_103	0.920	Male	Has relevent experience	no_enrollment	12th pass
1	29725	city_40	0.776	Male	No relevent experience	no_enrollment	12th pass
2	11561	city_21	0.624	NaN	No relevent experience	Full time course	12th pass
3	33241	city_115	0.789	NaN	No relevent experience	NaN	12th pass
4	666	city_162	0.767	Male	Has relevent experience	no_enrollment	12th pass

Data Quality Check

In [3]: `# Check shape of data`
`df.shape`

Out[3]: (21287, 14)

In [4]: `# Check for missing value percentage`
`(df.isnull().sum() / df.shape[0])*100`

Out[4]:

```

enrollee_id      0.000000
city              0.000000
city_development_index  0.000000
gender           23.563677
relevent_experience  0.000000
enrolled_university  1.958942
education_level    2.405224
major_discipline   14.680321
experience         0.328839
company_size       30.816931
company_type       31.822239
last_new_job       2.175036
training_hours     0.000000
target            10.001409
dtype: float64


```

In [5]: `# Drop rows where target value is missing`
`df.dropna(subset=['target'],inplace=True)`

```
# Convert target to int type
df['target'] = df['target'].astype('int')
```

```
In [6]: # Inspect missing values
msno.matrix(df)

plt.savefig('MissingValueMatrix.jpg',dpi=300, bbox_inches = 'tight');
```



Most missing values seem to come from 4 features: company_type, company_size, gender, and major_discipline

Because of the amount of missing columns, we will attempt a couple of methods to working with the missing data.

1. Use a Sklearn's SimpleImputer and replace missing values with 'most frequent'
2. Create an entirely new category called 'missing'

We will use these two methods for our pipeline.

Before imputing data, perform initial exploratory analysis.

Exploratory data analysis

```
In [7]: # Observe target distribution
pd.DataFrame(df.target.value_counts()).reset_index()
```

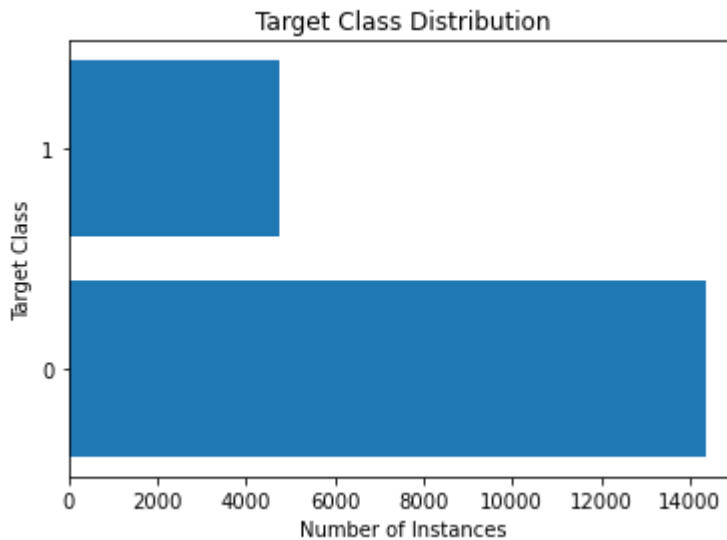
```
Out[7]:
```

	index	target
0	0	14381
1	1	4777

```
In [8]: # Plot target distribution
plot_target_bar = pd.DataFrame(df.target.value_counts()).reset_index()
x, y = plot_target_bar['index'].astype(str), plot_target_bar['target']
plt.ylabel('Target Class')
```

```
plt.xlabel('Number of Instances')
plt.title('Target Class Distribution')
plt.barh(x,y)

plt.savefig('TargetDistribution.jpg',dpi=300, bbox_inches = 'tight');
```



As we can see, there is large imbalance between our target variables with our 0 class making up as much as 75% of the total number of instances in the dataset. Therefore, when evaluating our models, we need to not only include accuracy as a metric, but also precision and recall.

```
In [9]: # Peek data
df.head()
```

```
Out[9]:
```

	enrollee_id	city	city_development_index	gender	relevent_experience	enrolled_university	educa
0	8949	city_103	0.920	Male	Has relevent experience	no_enrollment	
1	29725	city_40	0.776	Male	No relevent experience	no_enrollment	
2	11561	city_21	0.624	NaN	No relevent experience	Full time course	
3	33241	city_115	0.789	NaN	No relevent experience		NaN
4	666	city_162	0.767	Male	Has relevent experience	no_enrollment	

Observe for Skewed Data

```
In [10]: num_feats = df.dtypes[ df.dtypes!='object' ].index
skew_feats = df[num_feats].skew().sort_values(ascending=False)
skewness=pd.DataFrame({'Skew':skew_feats})
skewness
```

```
Out[10]:
```

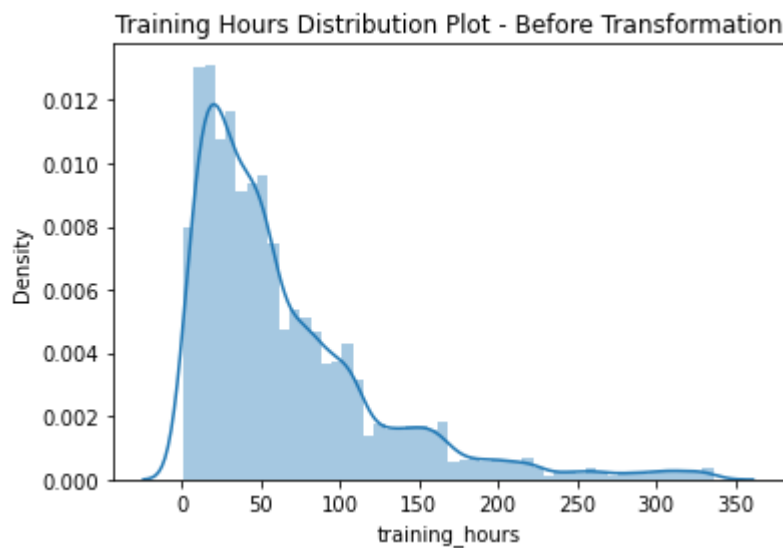
	Skew
--	------

	Skew
training_hours	1.819237
target	1.158815
enrollee_id	-0.018391
city_development_index	-0.995428

We will pay attention specifically to training_hours and city_development_index.

```
In [11]: training_hours = df['training_hours']
sns.distplot(training_hours)
plt.title('Training Hours Distribution Plot - Before Transformation')


plt.savefig('TrainingHours_BeforeTransformation.jpg',dpi=300, bbox_inches = 'tight');
```



There is an apparent right skew in training hours that we should address for our model. Let's observe how applying a log transformation could make the distribution more gaussian.

```
In [12]: log_training_hours = np.log(df['training_hours'])
sns.distplot(log_training_hours)
plt.title('Training Hours Distribution - Log Transformation')


plt.savefig('TrainingHours_AfterTransformation.jpg',dpi=300, bbox_inches = 'tight');
```



Applying a log transformation towards the training hours column appears to make the distribution more gaussian and suitable for modeling purposes.


```
In [13]: city_dev_index = df['city_development_index']
sns.distplot(city_dev_index)
plt.title('City Development Index Distribution Plot- Before Transformation')

plt.savefig('CityDevIndex_BeforeTransformation.jpg',dpi=300, bbox_inches = 'tight');
```



There appears to be a bimodal peaks with a left skew on the city_development_index column. Let's aim to apply a few transformation to address this in our pipeline.

```
In [14]: log_city_dev_index = np.log(df['city_development_index'])
sns.distplot(log_city_dev_index)
plt.title('Log City Development Index Distribution Plot');
```




The log transformation on the city_development_index column appeared to not have a major impact to making the distribution more gaussian. We will aim other transformation methods

```
In [15]: power = PowerTransformer(method='yeo-johnson', standardize=True)
pwr_city_dev = power.fit_transform( df[['city_development_index']] )
```

```
In [16]: sns.distplot(pwr_city_dev)
plt.title('City Development Index - Power Transformation')

plt.savefig('CityDevIndex_AfterTransformation.jpg',dpi=300, bbox_inches = 'tight');
```



A power transformation appears to work better in making the distribution more gaussian. The bimodal peak is still apparent.

```
In [17]: # Check for any apparent correlations -- specifically around city_development_index and
df[[x for x in df.columns if x != 'enrollee_id']].corr()
```

```
Out[17]:
```

	city_development_index	training_hours	target
city_development_index	1.000000	0.001920	-0.341665
training_hours	0.001920	1.000000	-0.021577
target	-0.341665	-0.021577	1.000000

We are able to observe there is no significant correlation between the numeric variables and our target. The feature with the highest correlation is city_development_index.

We observe an imbalance with our target. The imbalance is something we should address using stratified Kfold

```
In [18]: # Drop gender column to prevent against gender bias
df.drop(['enrollee_id', 'gender'], axis=1, inplace=True)

df.head()
```

```
Out[18]:
```

	city	city_development_index	relevent_experience	enrolled_university	education_level	major_disc
0	city_103	0.920	Has relevent experience	no_enrollment	Graduate	
1	city_40	0.776	No relevent experience	no_enrollment	Graduate	
2	city_21	0.624	No relevent experience	Full time course	Graduate	
3	city_115	0.789	No relevent experience	NaN	Graduate	Business D
4	city_162	0.767	Has relevent experience	no_enrollment	Masters	

Since the majority of the columns in the data are categorical, let us organize the columns into nominal and ordinal types to encode them appropriately.

Set Categorical Encoding

```
In [19]: # Access the columns of type object
df.loc[:, df.dtypes == object].head()
```

```
Out[19]:
```

	city	relevent_experience	enrolled_university	education_level	major_discipline	experience	comp
0	city_103	Has relevent experience	no_enrollment	Graduate	STEM	>20	
1	city_40	No relevent experience	no_enrollment	Graduate	STEM	15	
2	city_21	No relevent experience	Full time course	Graduate	STEM	5	
3	city_115	No relevent experience	NaN	Graduate	Business Degree	<1	
4	city_162	Has relevent experience	no_enrollment	Masters	STEM	>20	

```
In [20]: # Separate the columns into nominal and ordinal type columns
nominal = ['city', 'relevent_experience', 'enrolled_university', 'major_discipline', 'comp
```



```
ordinal = ['education_level', 'experience', 'company_size', 'last_new_job']
```

Ordinal Encode Mapping

Set mapping and ordering required for Sklearn's OrdinalEncoder to use in pipeline

Education Level

```
In [21]: # Peek at unique values of education_level and order them to feed into sklearn's OrdinalEncoder
df.education_level.unique()
```

```
Out[21]: array(['Graduate', 'Masters', 'High School', nan, 'Phd', 'Primary School'],
              dtype=object)
```

```
In [22]: # Set the ordering for education_level
edu_level_ord = [['missing', 'Primary School', 'High School', 'Graduate', 'Masters', 'Phd']]
```

Experience

```
In [23]: # Peek at unique values of experience and order them to feed into sklearn's OrdinalEncoder
df.experience.unique()
```

```
Out[23]: array(['>20', '15', '5', '<1', '11', '13', '7', '17', '2', '16', '1', '4',
               '10', '14', '18', '19', '12', '3', '6', '9', '8', '20', nan],
              dtype=object)
```

```
In [24]: # Assuming we replace missing values with the string 'missing'
experience_ord = [['missing',
                  '<1',
                  '1',
                  '2',
                  '3',
                  '4',
                  '5',
                  '6',
                  '7',
                  '8',
                  '9',
                  '10',
                  '11',
                  '12',
                  '13',
                  '14',
                  '15',
                  '16',
                  '17',
                  '18',
                  '19',
                  '20',
                  '>20']]
```

Company Size

```
In [25]: df.company_size.unique()
```

```
Out[25]: array([nan, '50-99', '<10', '10000+', '5000-9999', '1000-4999', '10/49',
               '100-500', '500-999'], dtype=object)
```

```
In [26]: # Assuming we replace missing values with the string 'missing'
company_size_ord = [['missing',
                    '<10',
                    '10/49',
                    '50-99',
                    '100-500',
                    '500-999',
                    '1000-4999',
                    '5000-9999',
                    '10000+']]
```

Last New Job

```
In [27]: df.last_new_job.unique()
```

```
Out[27]: array(['1', '>4', 'never', '4', '3', '2', nan], dtype=object)
```

```
In [28]: # Assuming we replace missing values with the string 'missing'
last_new_job_ord = [['missing',
                    'never',
                    '1',
                    '2',
                    '3',
                    '4',
                    '>4']]
```

Set up simple model using Logistic Regression

We will test our simple model against three dataframes using three different methods of imputation:

- Dropping missing values
- 'missing' value (create a new category called 'missing') imputation dataframe
- 'most frequent' value imputation

After setting up our three sets of data to run our simple LogisticRegression model we will perform the following:

1. Split data
2. Apply transformations
3. Train logistic regression model

Now we've created two dataframes with two different types of imputation methods. We will test both and observe which performs better on a baseline model.

Missing as a new category - Train Test Split

```
In [29]: df_baseline = df.copy()

# Create a df where missing values are filled with 'missing', and as a result creating
df_missing_baseline = df_baseline.fillna('missing')
# Separate target and predictors

# 'missing' new value imputation
```

```
y_missing = df_missing_baseline.target
X_missing = df_missing_baseline.drop('target',axis=1)
```

```
In [30]: # Separate into train and test

# 'missing' new value imputation
X_train_mis, X_test_mis, y_train_mis, y_test_mis = train_test_split(X_missing, y_mis
```

Most frequent imputation - Train Test Split

```
In [31]: # Create a df where most_frequent is used as the strategy to address missing values

# Set up frequent_imputer
frequent_imputer = SimpleImputer(strategy='most_frequent')

# Create a copy of our original df called df_frequent_baseline to help us organize how
df_frequent_baseline = df.copy()

# Most frequent imputation
y_freq = df_frequent_baseline.target
X_freq = df_frequent_baseline.drop('target',axis=1)
```

```
In [32]: # Separate into train and test

# Most frequent imputation
X_train_freq, X_test_freq, y_train_freq, y_test_freq = train_test_split(X_freq, y_freq,
```

Drop missing values - Train Test Split

```
In [33]: # Retain original dataframe and create a new one to drop null values
df_drop = df.copy()
df_drop.dropna(inplace=True)

y_drop = df_drop.target
X_drop = df_drop.drop('target',axis=1)

# Create train test split from our dropped null values dataframe
X_train_drop, X_test_drop, y_train_drop, y_test_drop = train_test_split(X_drop, y_drop,
```

Set up ColumnTransformer and pipeline

```
In [34]: # Set up helper function for retrieving the index of a column given column name
def get_column_index(df, col_name):
    return df.columns.get_loc(col_name)
```

For mixed types of numeric, nominal, and ordinal features, use sklearn's ColumnTransformer.

For the two numeric columns we will use:

- Log Transform for Training Hours
- Power Transform for City Development Index

```
In [35]: #For mixed types of numeric, nominal, and ordinal features, use sklearn's ColumnTransfo
log_transform = FunctionTransformer(np.log1p, validate=True)
```

```
t = [
    ("log_transform", FunctionTransformer(np.log1p, validate=True), [get_column_index(d
    ("pwr_transform", PowerTransformer(method='yeo-johnson', standardize=True), [get_co
    ("nominal", OneHotEncoder(handle_unknown='ignore'), [get_column_index(df_missing_ba
    ("edu_ord", OrdinalEncoder(categories=edu_level_ord), [get_column_index(df_missing_b
    ("exp_ord", OrdinalEncoder(categories=experience_ord), [get_column_index(df_missing
    ("comp_size_ord", OrdinalEncoder(categories=company_size_ord), [get_column_index(df
    ("new_job_ord", OrdinalEncoder(categories=last_new_job_ord), [get_column_index(df_m
]
```

Run Simple Model

Run simple model against the three dataframes:


- Dropping missing values
- 'missing' value (create a new category called 'missing') imputation dataframe
- 'most frequent' value imputation

```
In [36]: # define model
lr_model = LogisticRegression(solver='liblinear')
# define transform
ct = ColumnTransformer(transformers=t)
# define pipeline
pipeline = Pipeline(steps=[('t', ct), ('m', lr_model)])
# fit the pipeline on the transformed data
pipeline.fit(X_train_drop, y_train_drop.astype('int'))
# make predictions
yhat_drop = pipeline.predict(X_test_drop)
```

```
In [37]: # Check classification metrics
print(classification_report(y_test_drop, yhat_drop))
```

	precision	recall	f1-score	support
0	0.91	0.92	0.92	2319
1	0.60	0.55	0.58	476
accuracy			0.86	2795
macro avg	0.75	0.74	0.75	2795
weighted avg	0.86	0.86	0.86	2795

```
In [38]: # Plot confusion matrix
plot_confusion_matrix(pipeline, X_test_drop, y_test_drop,
                      cmap=plt.cm.Blues);
```



```
In [39]: # Plot and examine ROC AUC Curve
fig, ax = plt.subplots(figsize=(10,8))

# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test_drop))]

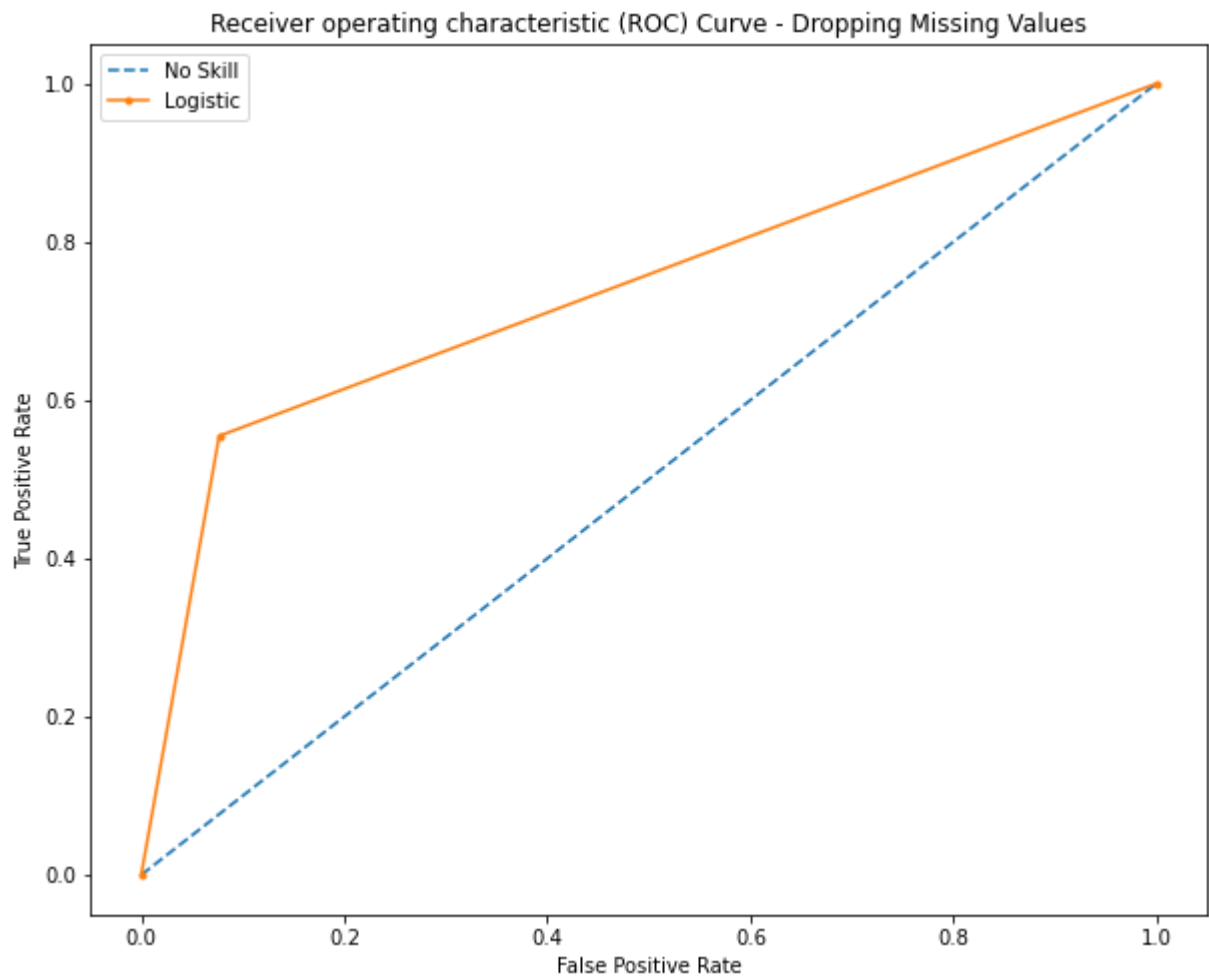
# calculate scores
ns_auc = roc_auc_score(y_test_drop, ns_probs)
lr_auc = roc_auc_score(y_test_drop, yhat_drop)
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('Logistic: ROC AUC=%.3f' % (lr_auc))
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test_drop, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test_drop, yhat_drop)
# plot the roc curve for the model
plt.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
plt.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

plt.title("Receiver operating characteristic (ROC) Curve - Dropping Missing Values")

# show the legend
plt.legend()

# Save fig
plt.savefig('LR_ROCCurve_drop.jpg',dpi=300, bbox_inches = 'tight');

No Skill: ROC AUC=0.500
Logistic: ROC AUC=0.739
```



Our simple logistic regression model against the dataframe with dropped missing values appears to perform decently with an AUC score of 0.7. However, let's continue with other methods of missing value imputation before moving forward.

Missing Value Simple Model Evaluation


```
In [40]: # define model
lr_model = LogisticRegression(solver='liblinear')
# define transform
ct = ColumnTransformer(transformers=t)
# define pipeline
pipeline = Pipeline(steps=[('t', ct), ('m', lr_model)])
# fit the pipeline on the transformed data
pipeline.fit(X_train_mis, y_train_mis.astype('int'))
# make predictions
yhat_mis = pipeline.predict(X_test_mis)
```

```
In [41]: # Check classification metrics
print(classification_report(y_test_mis, yhat_mis))
```

	precision	recall	f1-score	support
0	0.82	0.93	0.87	3621
1	0.61	0.35	0.44	1169
accuracy			0.79	4790
macro avg	0.71	0.64	0.66	4790

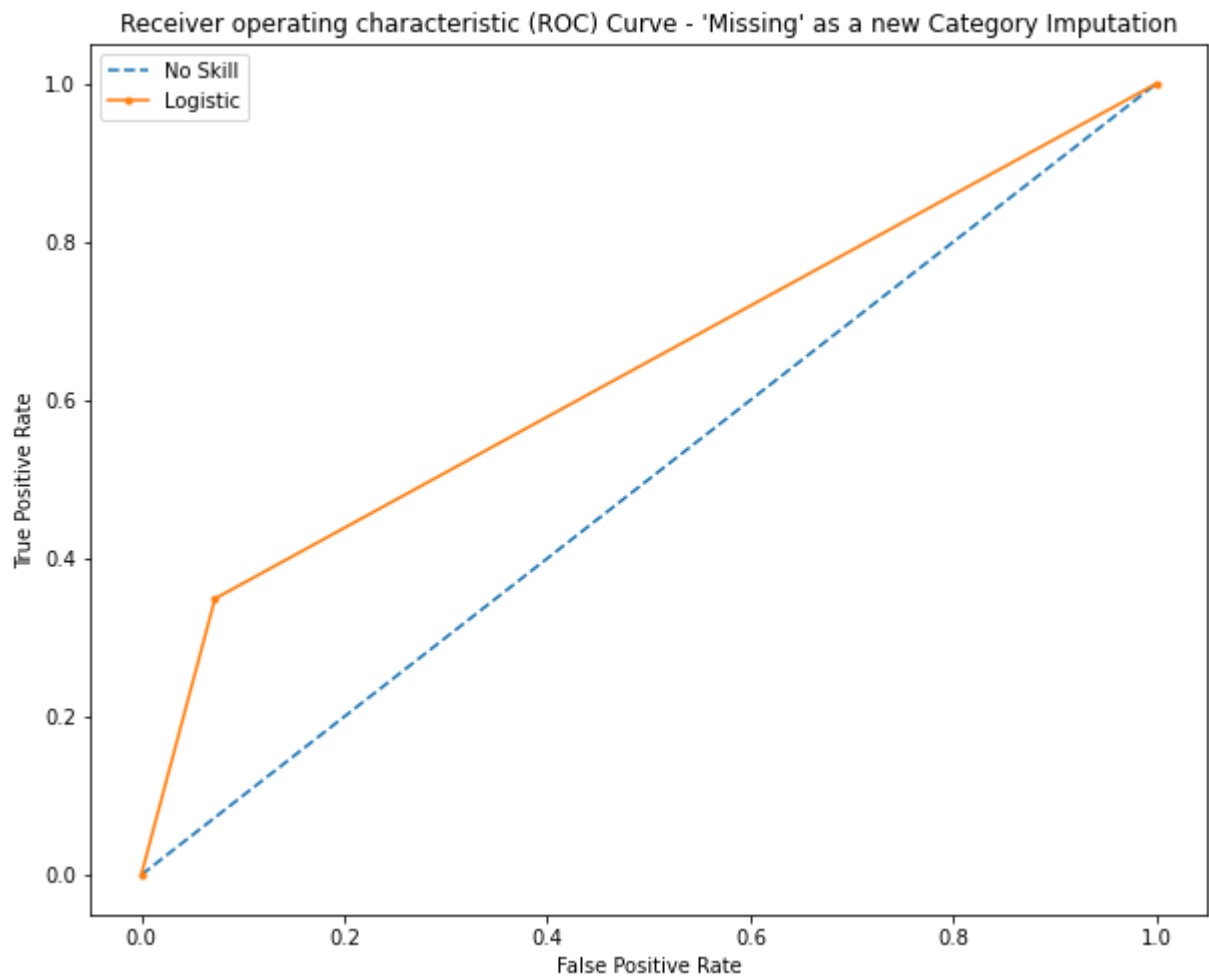
weighted avg 0.76 0.79 0.76 4790

```
In [42]: plot_confusion_matrix(pipeline, X_test_mis, y_test_mis,  
                               cmap=plt.cm.Blues);
```



```
In [43]: fig, ax = plt.subplots(figsize=(10,8))  
  
# generate a no skill prediction (majority class)  
ns_probs = [0 for _ in range(len(y_test_mis))]  
  
# calculate scores  
ns_auc = roc_auc_score(y_test_mis, ns_probs)  
lr_auc = roc_auc_score(y_test_mis, yhat_mis)  
# summarize scores  
print('No Skill: ROC AUC=%.3f' % (ns_auc))  
print('Logistic: ROC AUC=%.3f' % (lr_auc))  
# calculate roc curves  
ns_fpr, ns_tpr, _ = roc_curve(y_test_mis, ns_probs)  
lr_fpr, lr_tpr, _ = roc_curve(y_test_mis, yhat_mis)  
# plot the roc curve for the model  
plt.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')  
plt.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')  
# axis labels  
plt.xlabel('False Positive Rate')  
plt.ylabel('True Positive Rate')  
  
plt.title("Receiver operating characteristic (ROC) Curve - 'Missing' as a new Category")  
  
# show the Legend  
plt.legend()  
  
# Save fig  
plt.savefig('LR_ROCCurve_Missing.jpg',dpi=300, bbox_inches = 'tight');
```

No Skill: ROC AUC=0.500
Logistic: ROC AUC=0.638



Since 0 is our majority class making as much as 75% of our target class, it makes sense that its f1-score is significantly higher than the 1 class. In contrast, the 1 class has a poor f1 score. Specifically with low recall score of 0.34, this is our model is having a hard time classifying participants that will work for the data science company, actually will work for the company conducting data science training.

Most Frequent Imputation Simple Model Evaluation

```
In [44]: t = [
    ("log_transform", FunctionTransformer(np.log1p, validate=True), [get_column_index(df_frequent_b
    ("pwr_transform", PowerTransformer(method='yeo-johnson', standardize=True), [get_column_index(df_frequent_b
    ("nominal", OneHotEncoder(handle_unknown='ignore'), [get_column_index(df_frequent_b
    ("edu_ord", OrdinalEncoder(categories=edu_level_ord), [get_column_index(df_frequent_b
    ("exp_ord", OrdinalEncoder(categories=experience_ord), [get_column_index(df_frequent_b
    ("comp_size_ord", OrdinalEncoder(categories=company_size_ord), [get_column_index(df_frequent_b
    ("new_job_ord", OrdinalEncoder(categories=last_new_job_ord), [get_column_index(df_frequent_b
    ]
```

For our most frequent imputer imputation simple model evaluation, we want to set up a separate pipeline that includes the SimpleImputer object in a separate step. We need to include this as a separate step in the pipeline in order to prevent data leakage.

In the cell below, we will create a separate pipeline object called **pipeline_freq**

```
In [45]: # define model
```



```

lr_model = LogisticRegression(solver='liblinear')
# define transform
ct = ColumnTransformer(transformers=t)
# define pipeline
pipeline_freq = Pipeline(steps=[('imp', frequent_imputer), ('t', ct), ('m', lr_model)])
# fit the pipeline on the transformed data
pipeline_freq.fit(X_train_freq, y_train_freq.astype('int'))
# make predictions
yhat_freq = pipeline_freq.predict(X_test_freq)

```

```

In [46]: # Check classification metrics
y_test_freq = y_test_freq.astype('int')

print(classification_report(y_test_freq, yhat_freq))

```

	precision	recall	f1-score	support
0	0.81	0.92	0.86	3607
1	0.58	0.34	0.43	1183
accuracy			0.78	4790
macro avg	0.70	0.63	0.65	4790
weighted avg	0.75	0.78	0.76	4790

```

In [47]: fig, ax = plt.subplots(figsize=(10,8))

# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test_freq))]

# calculate scores
ns_auc = roc_auc_score(y_test_freq, ns_probs)
lr_auc = roc_auc_score(y_test_freq, yhat_freq)
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('Logistic: ROC AUC=%.3f' % (lr_auc))
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test_freq, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test_freq, yhat_freq)
# plot the roc curve for the model
plt.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
plt.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')


plt.title("Receiver operating characteristic (ROC) Curve - Most Frequent Imputation")

# show the legend
plt.legend()

# Save fig
plt.savefig('LR_ROCCurve_Freq.jpg', dpi=300, bbox_inches = 'tight');

```

No Skill: ROC AUC=0.500
 Logistic: ROC AUC=0.631



Both methods of imputation performed fairly similar both with the same accuracy of 77%. However, it appears the method simply dropping missing values appeared to perform best with an accuracy of 85% and an AUC score of 0.7.

Based on these initial results, we will move forward with simply dropping missing values.

Let's outline our next steps:

1. DecisionTreeClassifier
2. Extract feature importance from DecisionTreeClassifier
3. Experiment with RandomForest and XGBoost
4. Hyperparameter tuning

Loop over simple models and move forward with the best one

```
In [48]: # Encode our target to binary
y_drop = df_drop.target
X_drop = df_drop.drop('target',axis=1)

label_encoder = LabelEncoder()
y_drop = label_encoder.fit_transform(y_drop)
```

```

In [49]: # Instantiate models to loop over
dt = DecisionTreeClassifier()
xgb_clf = xgb.XGBClassifier()
lr = LogisticRegression(solver='liblinear')
rf = RandomForestClassifier()
svc = SVC()

# Create model list to iterate over
# models = [lr, dt, rf, xgb]

models = {'Logistic Regression': lr,
          'Support Vector Machine': svc,
          'Decision Tree Classifier': dt,
          'Random Forest': rf,
          'XGBoost': xgb_clf
        }

model_results = {'Model Name': [],
                 'Fold': [],
                 'Simple Model Accuracy Train': [],
                 'Simple Model Accuracy Test': []
                }

# Instantiate Stratified KFold
skf = StratifiedKFold(n_splits=5)

k_fold = 0

for train_index, test_index in tqdm(skf.split(X_drop, y_drop)):
    X_train, X_test = X_drop.iloc[train_index], X_drop.iloc[test_index]
    y_train, y_test = y_drop[train_index], y_drop[test_index]
    k_fold += 1
    for model_label, model in models.items():
        ct = ColumnTransformer(transformers=t)
        # define pipeline
        pipeline = Pipeline(steps=[('t', ct), ('m', model)])
        # fit the pipeline on the transformed data
        pipeline.fit(X_train, y_train.astype('int'))
        # make predictions
        yhat_drop = pipeline.predict(X_test)
        yhat_train_drop = pipeline.predict(X_train)
        # Calculate accuracy score
        score_train = accuracy_score(y_train, yhat_train_drop)
        score_train = round(score_train, 2)*100
        score_test = accuracy_score(y_test, yhat_drop)
        score_test = round(score_test, 2)*100
        # Append results
        model_results['Model Name'].append(model_label)
        model_results['Fold'].append(k_fold)
        model_results['Simple Model Accuracy Train'].append(score_train)
        model_results['Simple Model Accuracy Test'].append(score_test)

```

0it [00:00, ?it/s]

[20:24:03] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

1it [00:11, 11.29s/it]

[20:24:14] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the obj

ective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

2it [00:22, 11.34s/it]

[20:24:26] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

3it [00:34, 11.43s/it]

[20:24:37] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

4it [00:45, 11.39s/it]

[20:24:48] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.

5it [00:56, 11.29s/it]

```
In [50]: # Prepare and aggregate the accuracy scores from the model runs on the kfolds
results_df = pd.DataFrame( model_results )

results_agg_df = results_df.groupby('Model Name')[['Simple Model Accuracy Train', 'Simple Model Accuracy Test']]
results_agg_df
```

```
Out[50]:
```

	Model Name	Simple Model Accuracy Train	Simple Model Accuracy Test
0	Decision Tree Classifier	100.0	77.4
1	Logistic Regression	85.0	85.0
2	Random Forest	100.0	84.0
3	Support Vector Machine	84.6	84.6
4	XGBoost	90.4	84.2

Simple model candidate results

```
In [51]: def autolabel(rects):
        """
        Attach a text label above each bar displaying its height
        """
        for rect in rects:
            height = rect.get_height()
            ax.text(rect.get_x() + rect.get_width()/2., 1.01*height,
                    '%d' % int(height),
                    ha='center', va='bottom')
```

Code above is modified and borrowed from <https://stackoverflow.com/questions/45177937/how-can-i-adapt-the-autolabel-function-in-matplotlib-so-that-it-displays-negative-values>

```
In [52]: labels = results_agg_df['Model Name']
train_means = results_agg_df['Simple Model Accuracy Train']
test_means = results_agg_df['Simple Model Accuracy Test']
```

```

x = np.arange(len(labels)) # the Label Locations
width = 0.35 # the width of the bars

fig, ax = plt.subplots(figsize=(12,8))
rects1 = ax.bar(x - width/2, train_means, width, label='Train Sets')
rects2 = ax.bar(x + width/2, test_means, width, label='Test Sets')

# Add some text for Labels, title and custom x-axis tick labels, etc.
ax.set_ylabel('Accuracy Scores')
ax.set_title('Scores by Train and Test')
ax.set_xticks(x)
ax.set_xticklabels(labels)
ax.legend()


autolabel(rects1)
autolabel(rects2)

plt.xticks(rotation=45)

fig.tight_layout()

plt.savefig('ScoresbyTrainandTest.jpg',dpi=300, bbox_inches = 'tight')

```



Notably, there are signs of overfitting occurring within our DecisionTree, SVM, and XGBoost simple models given the higher train set accuracy scores. Nonetheless, it appears Logistic Regression has the highest average test accuracy score and so this is likely the model we want to move forward with and with an optimize.

Moreover, given the popularity, and popular hyperparameter optimization packages available for XGBoost, this is also a model I would like to test and move forward with despite the slight overfitting observed.

In terms of optimization, we will attempt the following:

- Extract feature importance from our more interpretable model -- Decision Tree Classifier
- Use Grid Search to find the optimal parameters for our SVM model
- Use the popular Optuna hyperparameter optimization library to improve our XGBoost simple model.

Feature Importance

```
In [53]: #For mixed types of numeric, nominal, and ordinal features, use sklearn's ColumnTransformer
log_transform = FunctionTransformer(np.log1p, validate=True)

t2 = [
    ("log_transform", FunctionTransformer(np.log1p, validate=True), [get_column_index(d
    ("pwr_transform", PowerTransformer(method='yeo-johnson', standardize=True), [get_co
    ("edu_ord", OrdinalEncoder(categories=edu_level_ord), [get_column_index(df_drop, 'ed
    ("exp_ord", OrdinalEncoder(categories=experience_ord), [get_column_index(df_drop, '
    ("comp_size_ord", OrdinalEncoder(categories=company_size_ord), [get_column_index(df
    ("new_job_ord", OrdinalEncoder(categories=last_new_job_ord), [get_column_index(df_d
]

# Columns transformed = ['training_hours', 'city_development_index', 'education_level', '
```

```
In [54]: # Encode our target to binary
y_drop = df_drop.target
X_drop = df_drop.drop('target', axis=1)

label_encoder = LabelEncoder()
y_drop = label_encoder.fit_transform(y_drop)
```

```
In [55]: ct2 = ColumnTransformer(transformers=t2)
X_drop_t = ct2.fit_transform(X_drop)
```

```
In [56]: # Separate the numeric features of our dataset
cols = ['training_hours', 'city_development_index', 'education_level', 'experience', 'comp
X_drop_t = pd.DataFrame(X_drop_t, columns = cols)

X_drop_t.head()
```

```
Out[56]:
```

	training_hours	city_development_index	education_level	experience	company_size	last_new_job
0	3.871201	-0.993870	3.0	16.0	3.0	6.0
1	2.197225	-1.052303	4.0	22.0	3.0	5.0
2	2.944439	-1.082999	3.0	14.0	1.0	6.0
3	3.850148	0.782225	3.0	8.0	3.0	2.0
4	4.820282	0.782225	3.0	18.0	8.0	6.0

```
In [57]: # Use pd.get_dummies to encode our nominal category columns
major_discipline_dummies = pd.get_dummies(X_drop.major_discipline)
major_discipline_dummies.columns = ["major_discipl_{i}".format(i) for i in major_discipl

company_type_dummies = pd.get_dummies(X_drop.company_type)
```

```
company_type_dummies.columns = ["company_type_{}".format(i) for i in company_type_dummies.columns]


# Concat the encoding of our major discipline and company type features
X_drop_dummies = pd.concat([major_discipline_dummies, company_type_dummies], axis=1)
```

```
In [58]: # Concatenate transformed columns
X_drop_fin = pd.concat([X_drop_t.reset_index(drop=True), X_drop_dummies.reset_index(drop=True)], axis=1)
X_drop_fin.fillna(0, inplace=True)
print(X_drop_fin.shape)
X_drop_fin.head()
```

(11179, 18)

```
Out[58]:
```

	training_hours	city_development_index	education_level	experience	company_size	last_new_job	major_discipline
0	3.871201	-0.993870	3.0	16.0	3.0	6.0	0
1	2.197225	-1.052303	4.0	22.0	3.0	5.0	0
2	2.944439	-1.082999	3.0	14.0	1.0	6.0	0
3	3.850148	0.782225	3.0	8.0	3.0	2.0	0
4	4.820282	0.782225	3.0	18.0	8.0	6.0	0




```
In [59]: # Instantiate and fit DecisionTreeClassifier
dt = DecisionTreeClassifier()
dt.fit(X_drop_fin, y_drop)
```

Out[59]: DecisionTreeClassifier()

```
In [60]: # Plot feature importances.
fig, ax = plt.subplots(figsize=(12,8))
ax.barh(X_drop_fin.columns, dt.feature_importances_, color='tab:blue')
ax.set(title='Feature Importances')
fig.tight_layout()

plt.savefig('FeatureImportances.jpg', dpi=300, bbox_inches = 'tight')
```



In order to get a better sense of the features that have a higher relevance in determining whether a participant will work for the data science training company or look for a new job, we used a simple DecisionTreeClassifier and extracted its determined feature importances.

As a result, we are able to observe that the city development index of the city where the participant is from plays the most important role in determining whether they will look for a new job or not. In descending order, the top 6 features include:

1. City Development Index (from where the participant is from)
2. Training Hours Completed
3. Years of Experience
4. Current Company Size
5. Difference in years between previous job and current job
6. Highest Level of Education

Hyperparameter Tuning - Logistic Regression

```
In [61]: # Encode our target to binary
y_drop = df_drop.target
X_drop = df_drop.drop('target',axis=1)

relevent_features = ['training_hours','city_development_index','education_level','exper
X_drop = X_drop[relevent_features]

label_encoder = LabelEncoder()
y_drop = label_encoder.fit_transform(y_drop)
```



```
In [62]: # Most dropout imputation and train test split
X_train_drop, X_test_drop, y_train_drop, y_test_drop = train_test_split(X_drop, y_drop,
```

```
In [63]: # Set up modified transformer for the relevant features selected from previous step
t3 = [
    ("log_transform", FunctionTransformer(np.log1p, validate=True), [get_column_index(X
    ("pwr_transform", PowerTransformer(method='yeo-johnson', standardize=True), [get_co
    ("edu_ord", OrdinalEncoder(categories=edu_level_ord), [get_column_index(X_train_drop
    ("exp_ord", OrdinalEncoder(categories=experience_ord), [get_column_index(X_train_dr
    ("comp_size_ord", OrdinalEncoder(categories=company_size_ord), [get_column_index(X_
    ("new_job_ord", OrdinalEncoder(categories=last_new_job_ord), [get_column_index(X_tr
]
```

```
In [64]: # Reuse previously created ColumnTransformer and Pipeline
ct = ColumnTransformer(transformers=t3)

# Instantiate our SVM
lr = LogisticRegression()

X_train_drop = ct.fit_transform(X_train_drop)
```

```
In [65]: # # defining parameter range
kfolds = StratifiedKFold(3)

param_grid = { 'solver': ['newton-cg', 'lbfgs', 'liblinear'],
               'penalty' : ['l1', 'l2'],
               'C' : [100, 10, 1.0, 0.1, 0.01] }

grid = GridSearchCV(lr, param_grid, refit = True, cv=kfolds.split(X_train_drop,y_train_

# fitting the model for grid search
grid.fit(X_train_drop, y_train_drop)
```

```
Out[65]: GridSearchCV(cv=<generator object _BaseKFold.split at 0x000001AF84EFCBA0>,
                    estimator=LogisticRegression(),
                    param_grid={'C': [100, 10, 1.0, 0.1, 0.01],
                                'penalty': ['l1', 'l2'],
                                'solver': ['newton-cg', 'lbfgs', 'liblinear']})
```

```
In [66]: # Examine best params
print(f'Best Params: {grid.best_params_}')

Best Params: {'C': 0.1, 'penalty': 'l1', 'solver': 'liblinear'}
```

```
In [67]: # Set up new LR model using grid.best_params_
lr_best = LogisticRegression(**grid.best_params_)
```

```
In [68]: # Fit the best params derived from the grid search towards the train set
lr_best.fit(X_train_drop, y_train_drop)

# Perform ColumnTransformer on Test set
X_test_drop = ct.fit_transform(X_test_drop)

# Perform predictions using best params from GridSearch
yhat_lr_best = lr_best.predict(X_test_drop)

print(classification_report(y_test_drop, yhat_lr_best))
```

```
precision    recall  f1-score   support
```

0	0.84	0.98	0.90	3043
1	0.57	0.14	0.23	647
accuracy			0.83	3690
macro avg	0.71	0.56	0.57	3690
weighted avg	0.79	0.83	0.79	3690

Plot ROC AUC Curve - LR with Hyperparameter Tuning

```
In [69]: fig, ax = plt.subplots(figsize=(10,8))

# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test_drop))]

# calculate scores
ns_auc = roc_auc_score(y_test_drop, ns_probs)
lr_auc = roc_auc_score(y_test_drop, yhat_lr_best)
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('Logistic: ROC AUC=%.3f' % (lr_auc))
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test_drop, ns_probs)
lr_fpr, lr_tpr, _ = roc_curve(y_test_drop, yhat_lr_best)
# plot the roc curve for the model
plt.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
plt.plot(lr_fpr, lr_tpr, marker='.', label='Logistic')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')


plt.title("Receiver operating characteristic (ROC) Curve - LR Hyperparameter Tuning")

# show the legend
plt.legend()

# Save fig
plt.savefig('LRFinal_ROCCurve_drop.jpg',dpi=300, bbox_inches = 'tight');
```

No Skill: ROC AUC=0.500

Logistic: ROC AUC=0.560



It appears that there is a slight drop in accuracy from our simple logistic regression model to our hypertuned model (85% --> 83%). Moreover, the true positive rate is still quite low, hovering slightly less than 0.2.

It is worth noting that while the model is performing well in predicting whether a person will work for the company (class 0), the model still struggles in understanding whether a person will look for a new job, and so it is still not very informative in telling the company whether a person will look for a new job (class 1), and hence hamper the company's ability to adjust to the characteristics and the course for folks looking to leave their job.

Hyperparameter Tuning - XGBoost

Code below is modified and borrowed from optuna.org

```
In [70]: def objective(trial):  
    y_drop = df_drop.target  
    X_drop = df_drop.drop('target', axis=1)  
  
    label_encoder = LabelEncoder()  
    y_drop = label_encoder.fit_transform(y_drop)  
  
    skf = StratifiedKFold(n_splits=5)  
  
    scores = []
```

```

for train_index, test_index in skf.split(X_drop,y_drop):
    train_x, test_x = X_drop.iloc[train_index], X_drop.iloc[test_index]
    train_y, test_y = y_drop[train_index], y_drop[test_index]

    ct = ColumnTransformer(transformers=t2)

    train_x = ct.fit_transform(train_x)
    test_x = ct.fit_transform(test_x)

    dtrain = xgb.DMatrix(train_x, label=train_y, enable_categorical=True)
    dtest = xgb.DMatrix(test_x, label=test_y, enable_categorical=True)

    param = {
#         "silent": 1,
        "objective": "binary:logistic",
        "eval_metric": "auc",
        "booster": trial.suggest_categorical("booster", ["gbtree", "gblinear", "dart"]),
        "lambda": trial.suggest_loguniform("lambda", 1e-8, 1.0),
        "alpha": trial.suggest_loguniform("alpha", 1e-8, 1.0),
    }

    if param["booster"] == "gbtree" or param["booster"] == "dart":
        param["max_depth"] = trial.suggest_int("max_depth", 1, 9)
        param["eta"] = trial.suggest_loguniform("eta", 1e-8, 1.0)
        param["gamma"] = trial.suggest_loguniform("gamma", 1e-8, 1.0)
        param["grow_policy"] = trial.suggest_categorical("grow_policy", ["depthwise", "lossguide"])
    if param["booster"] == "dart":
        param["sample_type"] = trial.suggest_categorical("sample_type", ["uniform", "poisson"])
        param["normalize_type"] = trial.suggest_categorical("normalize_type", ["tree", "loss"])
        param["rate_drop"] = trial.suggest_loguniform("rate_drop", 1e-8, 1.0)
        param["skip_drop"] = trial.suggest_loguniform("skip_drop", 1e-8, 1.0)

    bst = xgb.train(param, dtrain)
    preds = bst.predict(dtest)
    pred_labels = np.rint(preds)
    accuracy = accuracy_score(test_y, pred_labels)

    scores.append(accuracy)
return np.mean(scores)

```

```

In [71]: # Suppress Log statements
optuna.logging.set_verbosity(optuna.logging.WARNING)

# Run hyperparameter tuning study with optuna
study = optuna.create_study()
study.optimize(objective, n_trials=100, show_progress_bar=True)

```

```

In [72]: # Print best params
for key, value in study.best_trial.params.items():
    print(f"    {key}: {value}")

```

```

booster: dart
lambda: 6.25424868754657e-08
alpha: 0.3378532035423785
max_depth: 8
eta: 1.0845380910357793e-08
gamma: 0.0032795800919795575
grow_policy: depthwise
sample_type: uniform

```

```
normalize_type: forest
rate_drop: 2.5922928219214717e-05
skip_drop: 6.542375917105823e-07
```

```
In [87]: # Based on best params from optuna hyperparameter tuning, use the best params
params = {
```


```
    'booster': 'dart',
    'lambda': 6.25424868754657e-08,
    'alpha': 0.3378532035423785,
    'max_depth': 8,
    'eta': 1.0845380910357793e-08,
    'gamma': 0.0032795800919795575,
    'grow_policy': 'depthwise',
    'sample_type': 'uniform',
    'normalize_type': 'forest',
    'rate_drop': 2.5922928219214717e-05,
    'skip_drop': 6.542375917105823e-07
}
```

```
xgb_clf_fin = xgb.XGBClassifier(params=params)
```

```
In [88]: hist = study.trials_dataframe()
hist.head()
```

```
Out[88]:
```

	number	value	datetime_start	datetime_complete	duration	params_alpha	params_booste
0	0	0.832455	2021-07-15 20:24:54.415039	2021-07-15 20:24:55.128653	0 days 00:00:00.713614	9.805686e-05	gbtrei
1	1	0.833886	2021-07-15 20:24:55.131642	2021-07-15 20:24:55.812364	0 days 00:00:00.680722	1.468879e-08	dar
2	2	0.833886	2021-07-15 20:24:55.815357	2021-07-15 20:24:56.481579	0 days 00:00:00.666222	1.257218e-08	gbtrei
3	3	0.829323	2021-07-15 20:24:56.483572	2021-07-15 20:24:57.250641	0 days 00:00:00.767069	5.169953e-02	dar
4	4	0.833259	2021-07-15 20:24:57.253631	2021-07-15 20:24:57.961353	0 days 00:00:00.707722	3.229605e-02	gbtrei



```
In [89]: visualization.plot_optimization_history(study)
```

The plot above showcases the AUC score (y-axis) over the number of trials performed (x-axis). Through the 100 trials performed during hyperparameter tuning, it was observed that at trial 5, we received our highest AUC score of 0.834. From Optuna, we are able to pull the best params and perform final modeling.

Final Model with XGBoost Best Params

Set up Train Sample

```
In [77]: # Set up training
y = df_drop.target
X = df_drop.drop('target',axis=1)

# Encode target
label_encoder = LabelEncoder()
y = label_encoder.fit_transform(y)

# Train test split
# Use stratify parameter to preserve the target label proportions found in the original
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, stratify=y)

# Set Transformers
t_train = [
    ("log_transform", FunctionTransformer(np.log1p, validate=True), [get_column_index(X
    ("pwr_transform", PowerTransformer(method='yeo-johnson', standardize=True), [get_co
    ("nominal", OneHotEncoder(handle_unknown='ignore'), [get_column_index(X, x) for x i
    ("edu_ord", OrdinalEncoder(categories=edu_level_ord), [get_column_index(X, 'educatio
    ("exp_ord", OrdinalEncoder(categories=experience_ord), [get_column_index(X, 'experi
    ("comp_size_ord", OrdinalEncoder(categories=company_size_ord), [get_column_index(X,
    ("new_job_ord", OrdinalEncoder(categories=last_new_job_ord), [get_column_index(X, '
    ]

# Transform features
ct_train = ColumnTransformer(transformers=t_train)

# define pipeline
pipeline = Pipeline(steps=[('t', ct_train), ('m', xgb_clf_fin)])

# fit the pipeline on the transformed data
pipeline.fit(X_train, y_train)
```

```
[20:25:46] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:573:
Parameters: { "params" } might not be used.
```

This may not be accurate due to some parameters are only used in language bindings but passed down to XGBoost core. Or some parameters are not used but slip through this verification. Please open an issue if you find above cases.

```
[20:25:46] WARNING: C:/Users/Administrator/workspace/xgboost-win64_release_1.4.0/src/learner.cc:1095: Starting in XGBoost 1.3.0, the default evaluation metric used with the objective 'binary:logistic' was changed from 'error' to 'logloss'. Explicitly set eval_metric if you'd like to restore the old behavior.
```

```
Out[77]: Pipeline(steps=[('t',
                           ColumnTransformer(transformers=[('log_transform',
                                                              FunctionTransformer(func=<ufunc 'log1
p'>,
                                                              validate=True),
                                                              [10])),
                           ('pwr_transform',
                             PowerTransformer(), [1])),
                           ('nominal',
                             OneHotEncoder(handle_unknown='ignor
e'),
                           [0, 2, 3, 5, 8]),
                           ('edu_ord',
                             OrdinalEncoder(categories=[['missing',
                                                           'Primary',
                                                           'School',
                                                           'High',
                                                           'School',
                                                           'Graduat
e',
                                                           'Masters',
                                                           'Phd']]),
                           [4])...
                           learning_rate=0.300000012, max_delta_step=0,
                           max_depth=6, min_child_weight=1, missing=nan,
                           monotone_constraints='()', n_estimators=100,
                           n_jobs=8, num_parallel_tree=1,
                           params={'alpha': 0.10383266486970251,
                                   'booster': 'gbtree',
                                   'grow_policy': 'depthwise',
                                   'max_depth': 2},
                           random_state=0, reg_alpha=0, reg_lambda=1,
                           scale_pos_weight=1, subsample=1,
                           tree_method='exact', validate_parameters=1,
                           verbosity=None))])
```

Predict on Test

```
In [78]: preds = pipeline.predict(X_test)
```

```
In [79]: # Check classification metrics
print(classification_report(y_test, preds))
```

	precision	recall	f1-score	support
0	0.88	0.93	0.90	2293
1	0.56	0.39	0.46	502
accuracy			0.84	2795

macro avg	0.72	0.66	0.68	2795
weighted avg	0.82	0.84	0.82	2795

Plot ROC AUC Curve - XGB with Hyperparameter Tuning

```
In [80]: fig, ax = plt.subplots(figsize=(10,8))

# generate a no skill prediction (majority class)
ns_probs = [0 for _ in range(len(y_test))]

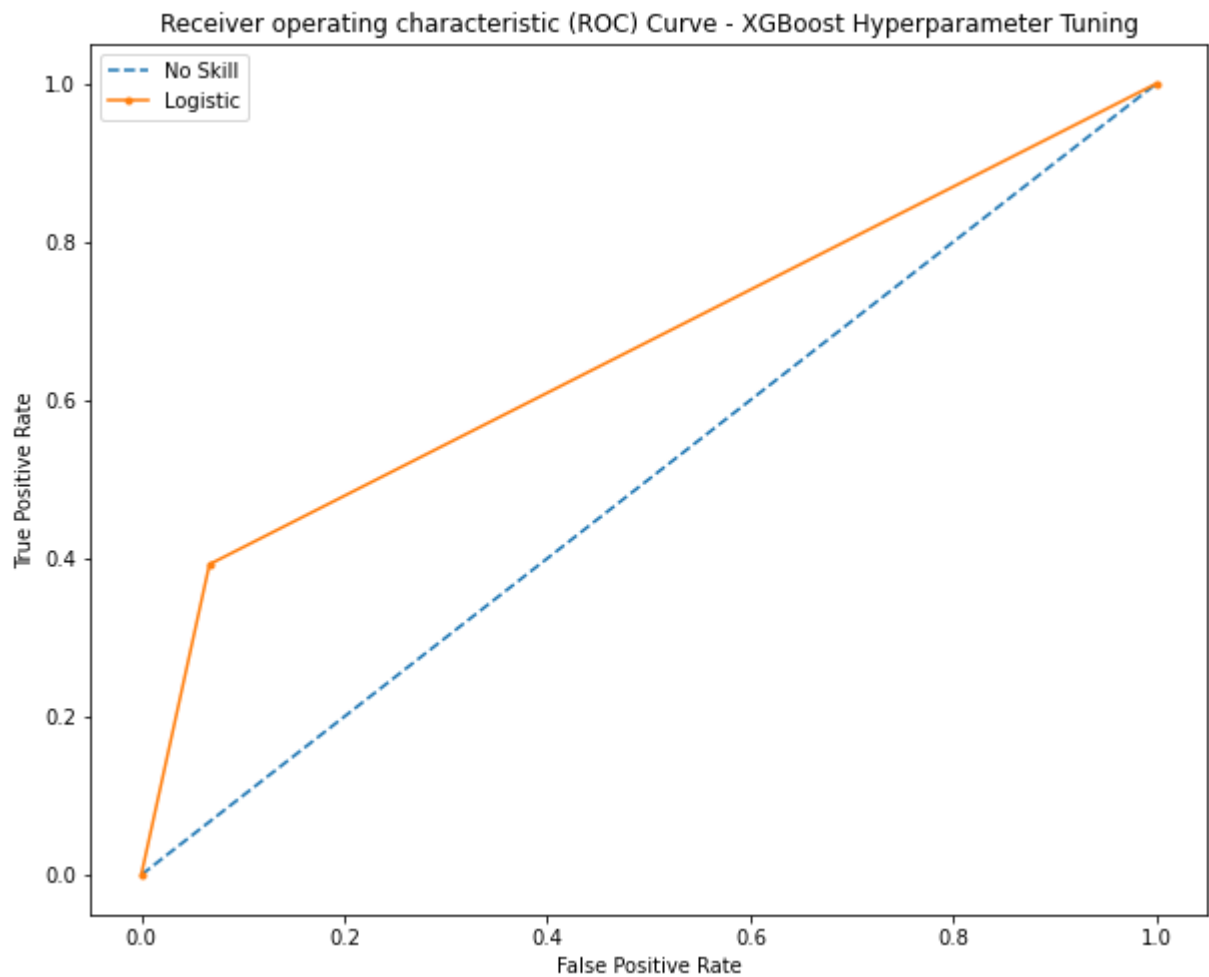
# calculate scores
ns_auc = roc_auc_score(y_test, ns_probs)
xgb_auc = roc_auc_score(y_test, preds)
# summarize scores
print('No Skill: ROC AUC=%.3f' % (ns_auc))
print('Logistic: ROC AUC=%.3f' % (xgb_auc))
# calculate roc curves
ns_fpr, ns_tpr, _ = roc_curve(y_test, ns_probs)
xgb_fpr, xgb_tpr, _ = roc_curve(y_test, preds)
# plot the roc curve for the model
plt.plot(ns_fpr, ns_tpr, linestyle='--', label='No Skill')
plt.plot(xgb_fpr, xgb_tpr, marker='.', label='Logistic')
# axis labels
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')

plt.title("Receiver operating characteristic (ROC) Curve - XGBoost Hyperparameter Tuning")

# show the legend
plt.legend()

# Save fig
plt.savefig('XGBFinal_ROCCurve_drop.jpg',dpi=300, bbox_inches = 'tight');
```

No Skill: ROC AUC=0.500
Logistic: ROC AUC=0.663



The tuned XGboost model performs much better than our tuned LR model (0.678 compared to 0.558). Therefore for future iterations of the model, we may wish to continue with XGBoost.

However, it is still worth noting that the model performance on predicting whether a person will leave the company (class 1) is poor with an f1-score of 0.49 compared to class 0's f1-score of 0.9. In order to improve this outcome, we will need more instances of class 1 to train against, and hence better understand the factor that lead a person to leave their job after data science training.

Characteristics of People Looking for a new Job

City Development Index Comparison

```
In [81]: fig, ax = plt.subplots(figsize=(10,8))


# Iterate through target label
for t in df_drop.target.unique():
    # Create temp dataframe for the genre being iterated on
    frame = df_drop[df_drop.target==t]

    # Create list of sample_means of size 100 to plot distributions of averagerating.
    # Set up city dev index comparison
    sample_means = [np.random.choice(frame.city_development_index, replace=True, size=1
```

```
plt.hist(sample_means, label=t,alpha=0.6)

plt.legend()
plt.ylabel('Count',fontsize=16)
plt.xlabel('City Development Index ',fontsize=16)
plt.title('City Development Index Target Comparison',fontsize=16)

# Format and save fig
plt.savefig('CityDevelopmentComparison.jpg',dpi=300, bbox_inches = 'tight')
```



In the visual above, I bootstrapped the target variables to generate a distribution over each targets' city development index to provide a comparison of how city development index plays as a factor for whether a person looks for a new job or not after completing data science training. We can observe that people who are not look for a job typically live in better developed cities compared to people in looking for jobs in lower developed cities.

Training Hours Comparison

```
In [82]: fig, ax = plt.subplots(figsize=(10,8))

# Iterate through target label
for t in df_drop.target.unique():
    # Create temp dataframe for the genre being iterated on
    frame = df_drop[df_drop.target==t]

    # Create list of sample_means of size 100 to plot distributions of averagerating.
```

```
# Set up city dev index comparison
sample_means = [np.random.choice(frame.training_hours, replace=True, size=len(frame))]
plt.hist(sample_means, label=t,alpha=0.6)

plt.legend()
plt.ylabel('Count',fontsize=16)
plt.xlabel('Training Hours',fontsize=16)
plt.title('Training Hours Target Comparison',fontsize=16)

# Format and save fig
plt.savefig('TrainingHours_Comparison.jpg',dpi=300, bbox_inches = 'tight')
```



In the visual above, I bootstrapped the target variables to generate a distribution over each targets' training hours to provide a comparison of how the number of training hours completed plays as a factor for whether a person looks for a new job or not after completing data science training. We can observe that people who are not looking for a new job typically spend more hours training versus people who are looking for a new job.

Experience Comparison

```
In [83]: # Collect 100 samples of the dataframe 100 times
plot_df = pd.DataFrame()
for i in tqdm(range(100)):
    experience_1_sample = df_drop[df_drop['target']==1].sample(n=1000,replace=True)
    experience_0_sample = df_drop[df_drop['target']==0].sample(n=1000,replace=True)
```

```

append_0 = pd.DataFrame(experience_0_sample.groupby(['target', 'experience']).size())
append_1 = pd.DataFrame(experience_1_sample.groupby(['target', 'experience']).size())

plot_df = plot_df.append(append_0)
plot_df = plot_df.append(append_1)

```

100%|██████████| 100/100 [00:01<00:00, 86.56it/s]

```

In [84]: print(plot_df.shape)
         plot_df.head()

```

(4397, 3)

```

Out[84]:

```

	target	experience	Count
0	0	1	10
1	0	10	59
2	0	11	45
3	0	12	34
4	0	13	23

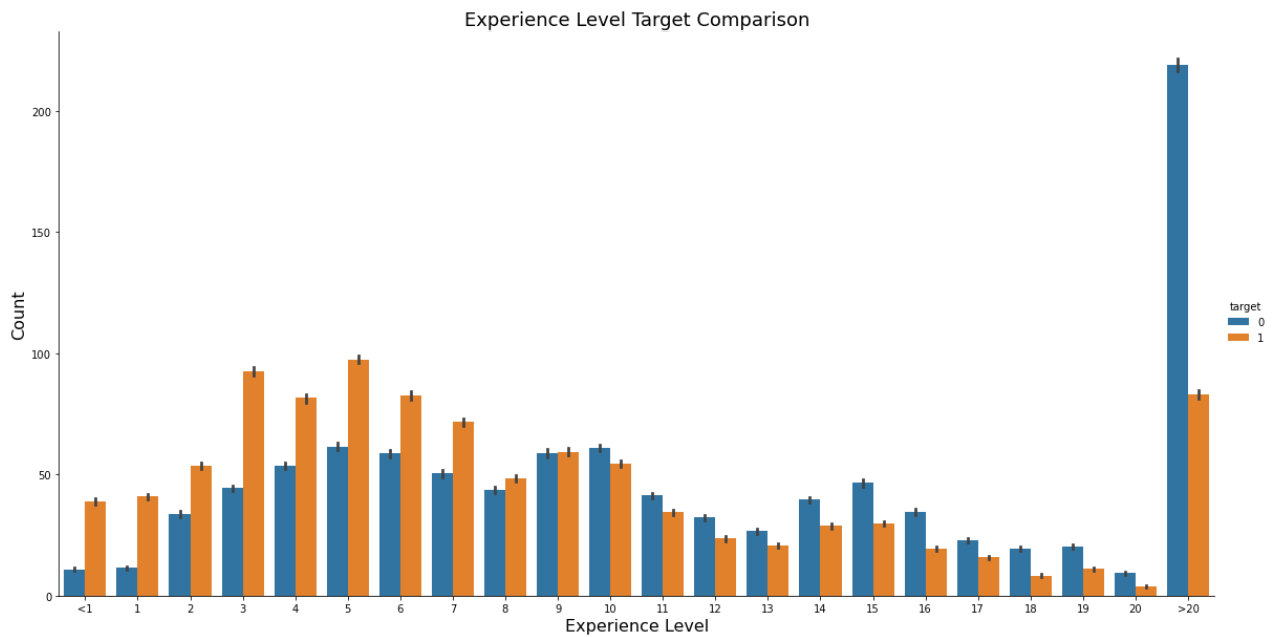
```

In [85]: # Plot df
         order = [ '<1',
                   '1',
                   '2',
                   '3',
                   '4',
                   '5',
                   '6',
                   '7',
                   '8',
                   '9',
                   '10',
                   '11',
                   '12',
                   '13',
                   '14',
                   '15',
                   '16',
                   '17',
                   '18',
                   '19',
                   '20',
                   '>20' ]

         sns.catplot(x="experience", y="Count", hue="target", kind="bar", data=plot_df, order=order)
         plt.xlabel('Experience Level', fontsize=16)
         plt.ylabel('Count', fontsize=16)
         plt.title('Experience Level Target Comparison', fontsize=18)

         # Format and save fig
         plt.savefig('ExperienceLevel_Comparison.jpg', dpi=300, bbox_inches = 'tight')

```



Lastly, through a similar bootstrapping method used in the above two visuals, I looked at how years of experience play as a factor for people looking for a new job or not. It appears that typically people with 9 years or less experience are more open to looking for new roles compared to people who have 10+ years of experience, where they are more likely to stay put in their current role.

Conclusion

Regarding the features that have the most impact of a person's decision into looking for a new role or not, I would highlight City Development Index (city development score of where the employee is from), the number of training hours completed, and the amount of experience an employee has as the top 3 factors. Moreover, from a modeling stand point, I would continue to iterate on the XGBoost model in hopes of collecting more data on people who are looking for a new job to counteract the imbalance of the dataset. Currently, our model is still performing quite poorly on our recall metric -- roughly 0.49 for our class of employees looking for a new role. Given the company's objective to reduce cost, and lost time for employees looking for a new role, we would want to correctly identify employees looking for a new role, and avoid mistaking these employees for those looking to stay. Otherwise, if we think an employee is staying, but in reality they are leaving, then there associated cost and time.

Next steps

Based on the presented analysis, there are more steps we can take to improve our model. One step is to continue experimenting with other methods of missing data imputation, and build iterative modelling on top of these other methods. Another step that will take a collective effort from the company is prioritizing data quality, and ensuring no further missing data. The last note to best improve the model is to collect more high quality data, especially data on employees looking to leave their role.

Sources

- "ROC Curves and Precision Recall Curves for Classification in Python." Machine Learning Mastery, <https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classification-in-python/>
- "How to Use ROC Curves and Precision-Recall Curves for Classification in Python", Machine Learning Mastery, <https://machinelearningmastery.com/roc-curves-and-precision-recall-curves-for-classification-in-python/>
- Brownlee, Jason. Data Preparation for Machine Learning: Data Cleaning, Feature Selection, and Data Transforms in Python. V 1.2 Data Preparation for Machine Learning, 2021